

hw2

January 22, 2024

1 HW2

Kevin Smith & Evelina Teran

```
[ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import Ridge
from sklearn.metrics import r2_score, mean_squared_error
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor

abalone_data = pd.read_csv('/Users/kevin_smith/Desktop/FSU_Relevant_Stuff/
↳spring23/STA5635/homework/hw2/abalone.csv')
```

```
[ ]: # Split data into predictors and response variables
X = abalone_data.iloc[:, :7]
y = abalone_data.iloc[:, -1]

# Number of random splits
num_splits = 20

## Code for Question 1
# Create a list to store the MSE values for training and testing sets
train_mse_list = []
test_mse_list = []

for _ in range(num_splits):
    # Random split
    X_train, X_test, y_train, y_test = train_test_split(X,
                                                         y,
                                                         test_size = 0.1,
                                                         random_state = np.random.randint(1,100))

    # Calculate average training y
    avg_train_y = np.mean(y_train)
```

```

# Predict test set responses using average training y
y_pred_train = np.full_like(y_train, avg_train_y)
y_pred_test = np.full_like(y_test, avg_train_y)

# Calculate MSE
train_mse = mean_squared_error(y_train, y_pred_train)
test_mse = mean_squared_error(y_test, y_pred_test)

# Append MSE values to lists
train_mse_list.append(train_mse)
test_mse_list.append(test_mse)

# Calculate average MSE for training and testing sets across the 20 splits
avg_null_train_mse = np.mean(train_mse_list)
avg_null_test_mse = np.mean(test_mse_list)

print(f"Average training MSE for Null Model: {avg_null_train_mse}")
print(f"Average testing MSE for Null Model: {avg_null_test_mse}")

```

Average training MSE for Null Model: 11.206506120276742

Average testing MSE for Null Model: 11.727153110047848

1.1 Question 2

```

[ ]: # Create a list to store metrics
train_r2_list, test_r2_list = [], []
train_mse_list, test_mse_list = [], []
log_det_list = []

for _ in range(num_splits):
    # Random split
    X_train, X_test, y_train, y_test = train_test_split(X,
                                                         y,
                                                         test_size = 0.1,
                                                         random_state = np.random.randint(1,100))

    # Perform Ridge Regression
    lambda_value = 0.001
    XTX_plus_lambdaIp = np.dot(X_train.T, X_train) + lambda_value * np.
    ↪identity(X_train.shape[1])

    ridge_weights = np.linalg.solve(XTX_plus_lambdaIp, np.dot(X_train.T, y_train))

    # Model Evaluation
    y_pred_train = np.dot(X_train, ridge_weights)
    y_pred_test = np.dot(X_test, ridge_weights)

```

```

train_r2 = r2_score(y_train, y_pred_train)
test_r2 = r2_score(y_test, y_pred_test)

train_mse = mean_squared_error(y_train, y_pred_train)
test_mse = mean_squared_error(y_test, y_pred_test)

# Calculate Log Determinant
log_det = np.log(np.linalg.det(XTX_plus_lambdaIp))

# Append metrics to lists
train_r2_list.append(train_r2)
test_r2_list.append(test_r2)
train_mse_list.append(train_mse)
test_mse_list.append(test_mse)
log_det_list.append(log_det)

# Calculate average and standard deviation for metrics
avg_train_r2 = np.mean(train_r2_list)
std_train_r2 = np.std(train_r2_list)

avg_test_r2 = np.mean(test_r2_list)
std_test_r2 = np.std(test_r2_list)

avg_train_mse = np.mean(train_mse_list)
std_train_mse = np.std(train_mse_list)

avg_test_mse = np.mean(test_mse_list)
std_test_mse = np.std(test_mse_list)

avg_log_det = np.mean(log_det_list)
std_log_det = np.std(log_det_list)

# Print results
print(f"Average Training R^2: {avg_train_r2}, Std Training R^2: {std_train_r2}")
print(f"Average Testing R^2: {avg_test_r2}, Std Testing R^2: {std_test_r2}")
print(f"Average Training MSE: {avg_train_mse}, Std Training MSE: {std_train_mse}")
print(f"Average Testing MSE: {avg_test_mse}, Std Testing MSE: {std_test_mse}")
print(f"Average Log Determinant: {avg_log_det}, Std Log Determinant: {std_log_det}")

```

Average Training R²: 0.5125645787609263, Std Training R²: 0.004308612628212074
 Average Testing R²: 0.5238960466683833, Std Testing R²: 0.0403581722774708
 Average Training MSE: 5.050367976968313, Std Training MSE: 0.07288960676216223
 Average Testing MSE: 5.023363245432512, Std Testing MSE: 0.6642493605933701
 Average Log Determinant: 18.290646036857108, Std Log Determinant:

0.06540525467986825

1.2 Question 3

```
[ ]: max_depths = range(1,8)

# List to store metrics
train_r2_avg, test_r2_avg = [], []
train_mse_avg, test_mse_avg = [], []

null_model_mse = avg_null_test_mse

for depth in max_depths:
    train_r2_list, test_r2_list = [], []
    train_mse_list, test_mse_list = [], []

    for _ in range(num_splits):
        # Random Split
        X_train, X_test, y_train, y_test = train_test_split(X,
                                                            y,
                                                            test_size=0.1,
                                                            random_state=np.random.randint(1,100))

        # Decision Tree Regression
        dt_model = DecisionTreeRegressor(max_depth=depth)
        dt_model.fit(X_train, y_train)

        # Model Evaluation
        y_pred_train = dt_model.predict(X_train)
        y_pred_test = dt_model.predict(X_test)

        train_r2 = r2_score(y_train, y_pred_train)
        test_r2 = r2_score(y_test, y_pred_test)

        train_mse = mean_squared_error(y_train, y_pred_train)
        test_mse = mean_squared_error(y_test, y_pred_test)

        # Append metrics
        train_r2_list.append(train_r2)
        test_r2_list.append(test_r2)
        train_mse_list.append(train_mse)
        test_mse_list.append(test_mse)

    # Calculate average metric for current tree depth
    avg_train_r2 = np.mean(train_r2_list)
    avg_test_r2 = np.mean(test_r2_list)
    avg_train_mse = np.mean(train_mse_list)
```

```

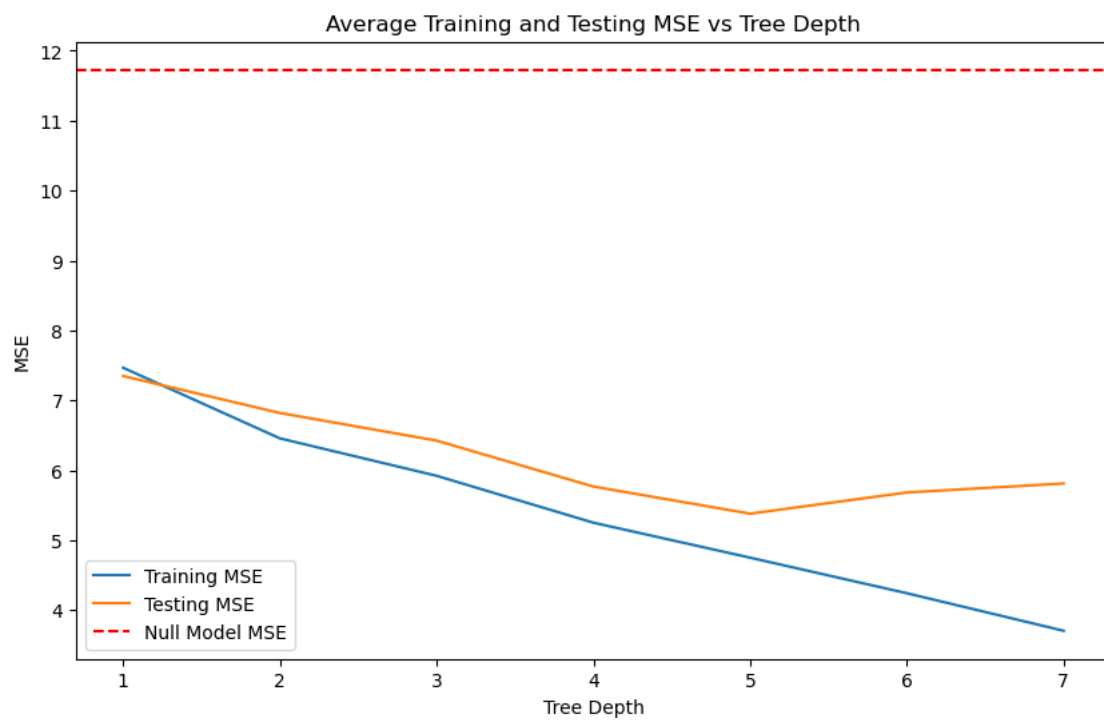
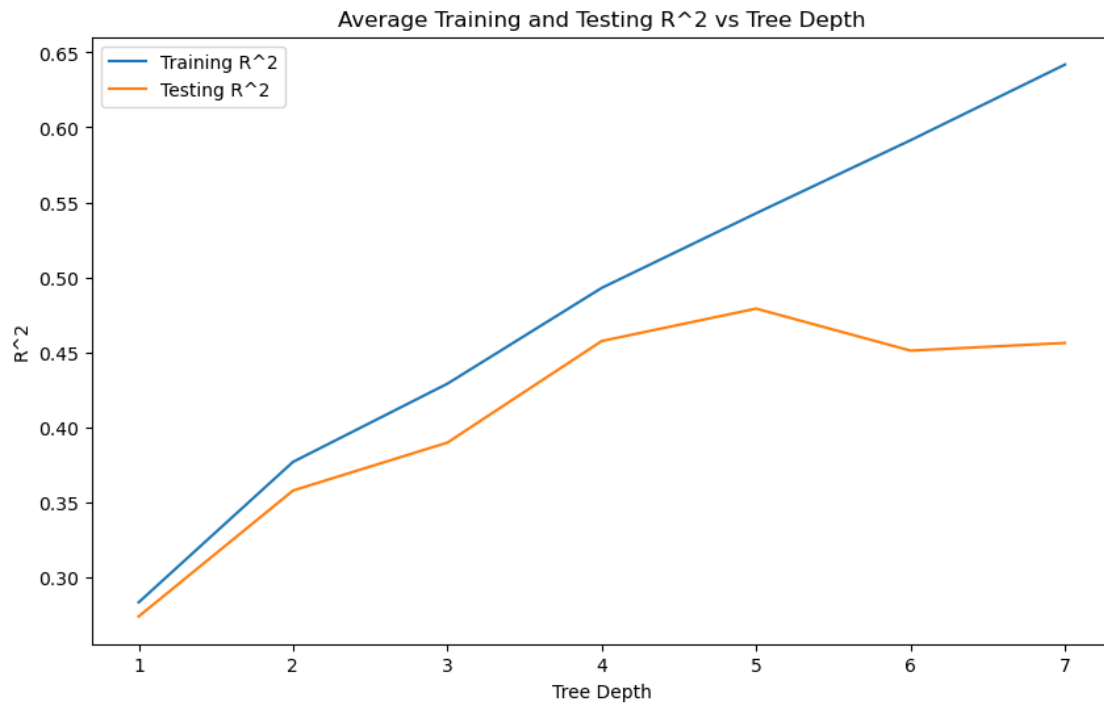
avg_test_mse = np.mean(test_mse_list)

# Append average metrics to lists
train_r2_avg.append(avg_train_r2)
test_r2_avg.append(avg_test_r2)
train_mse_avg.append(avg_train_mse)
test_mse_avg.append(avg_test_mse)

# Plot R^2 vs Tree Depth
plt.figure(figsize= (10,6))
plt.plot(max_depths, train_r2_avg, label="Training R^2")
plt.plot(max_depths, test_r2_avg, label="Testing R^2")
plt.xlabel("Tree Depth")
plt.ylabel("R^2")
plt.title("Average Training and Testing R^2 vs Tree Depth")
plt.legend()
plt.savefig("r2_vs_depth.png")

# Plot MSE vs Tree Depth with Null Model MSE as a horizontal line
plt.figure(figsize= (10,6))
plt.plot(max_depths, train_mse_avg, label="Training MSE")
plt.plot(max_depths, test_mse_avg, label="Testing MSE")
plt.axhline(y=null_model_mse, color='r', linestyle='--', label="Null Model MSE")
plt.xlabel("Tree Depth")
plt.ylabel("MSE")
plt.title("Average Training and Testing MSE vs Tree Depth")
plt.legend()
plt.savefig("MSE_vs_depth.png")

```



1.3 Question 4

```
[ ]: num_trees_list = [10, 30, 100, 300]

# Lists to store results
results = {}

for num_trees in num_trees_list:
    train_r2_list, test_r2_list = [], []
    train_mse_list, test_mse_list = [], []

    for _ in range(num_splits):
        # Random Split
        X_train, X_test, y_train, y_test = train_test_split(X,
                                                            y,
                                                            test_size=0.1,
                                                            random_state = np.random.randint(1, 100))

        # Random Forest Regression
        rf_model = RandomForestRegressor(n_estimators=num_trees)
        rf_model.fit(X_train, y_train)

        # Model Evaluation
        y_pred_train = rf_model.predict(X_train)
        y_pred_test = rf_model.predict(X_test)

        train_r2 = r2_score(y_train, y_pred_train)
        test_r2 = r2_score(y_test, y_pred_test)

        train_mse = mean_squared_error(y_train, y_pred_train)
        test_mse = mean_squared_error(y_test, y_pred_test)

        # Append metrics
        train_r2_list.append(train_r2)
        test_r2_list.append(test_r2)
        train_mse_list.append(train_mse)
        test_mse_list.append(test_mse)

    # Calculate average and std for each metric
    avg_train_r2 = np.mean(train_r2_list)
    std_train_r2 = np.std(train_r2_list)

    avg_test_r2 = np.mean(test_r2_list)
    std_test_r2 = np.std(test_r2_list)

    avg_train_mse = np.mean(train_mse_list)
    std_train_mse = np.std(train_mse_list)
```

```

avg_test_mse = np.mean(test_mse_list)
std_test_mse = np.std(test_mse_list)

# Store results in dictionary
results[num_trees] = {
    'avg_train_r2': avg_train_r2,
    'std_train_r2': std_train_r2,
    'avg_test_r2': avg_test_r2,
    'std_test_r2': std_test_r2,
    'avg_train_mse': avg_train_mse,
    'std_train_mse': std_train_mse,
    'avg_test_mse': avg_test_mse,
    'std_test_mse': std_test_mse
}

# Print results
for num_trees, metrics in results.items():
    print(f"Number of Trees: {num_trees}")
    print(f"Average Training R^2: {metrics['avg_train_r2']}, Std Training R^2:␣
↪{metrics['std_train_r2']}")
    print(f"Average Testing R^2: {metrics['avg_test_r2']}, Std Testing R^2:␣
↪{metrics['std_test_r2']}")
    print(f"Average Training MSE: {metrics['avg_train_mse']}, Std Training MSE:␣
↪{metrics['std_train_mse']}")
    print(f"Average Testing MSE: {metrics['avg_test_mse']}, Std Testing MSE:␣
↪{metrics['std_test_mse']}")
    print("\n")

```

Number of Trees: 10
 Average Training R²: 0.9121517840732917, Std Training R²:
 0.0021863033839376067
 Average Testing R²: 0.4937228547204806, Std Testing R²: 0.04341150846213091
 Average Training MSE: 0.9104748536455561, Std Training MSE: 0.023812358769338838
 Average Testing MSE: 5.351228468899522, Std Testing MSE: 0.4561297826041411

Number of Trees: 30
 Average Training R²: 0.9289959752413977, Std Training R²:
 0.0014472438105444635
 Average Testing R²: 0.5460791025735864, Std Testing R²: 0.04113804224999156
 Average Training MSE: 0.7369183223937081, Std Training MSE: 0.015222734220515241
 Average Testing MSE: 4.72877711323764, Std Testing MSE: 0.44132506032382796

Number of Trees: 100
 Average Training R²: 0.9353138910462484, Std Training R²: 0.00090008092177222

Average Testing R^2 : 0.5439984753433376, Std Testing R^2 : 0.03637546319873581
Average Training MSE: 0.6738444704630122, Std Training MSE: 0.01128145557897577
Average Testing MSE: 4.579415789473684, Std Testing MSE: 0.462465681183874

Number of Trees: 300

Average Training R^2 : 0.937485677101028, Std Training R^2 : 0.0006817558829635172
Average Testing R^2 : 0.539299799733937, Std Testing R^2 : 0.03501326062011526
Average Training MSE: 0.6505462144462184, Std Training MSE: 0.010185497945396565
Average Testing MSE: 4.711116133705476, Std Testing MSE: 0.4946323567039948